

NoSQLBench: Why ?



How/How-well do you answer these-

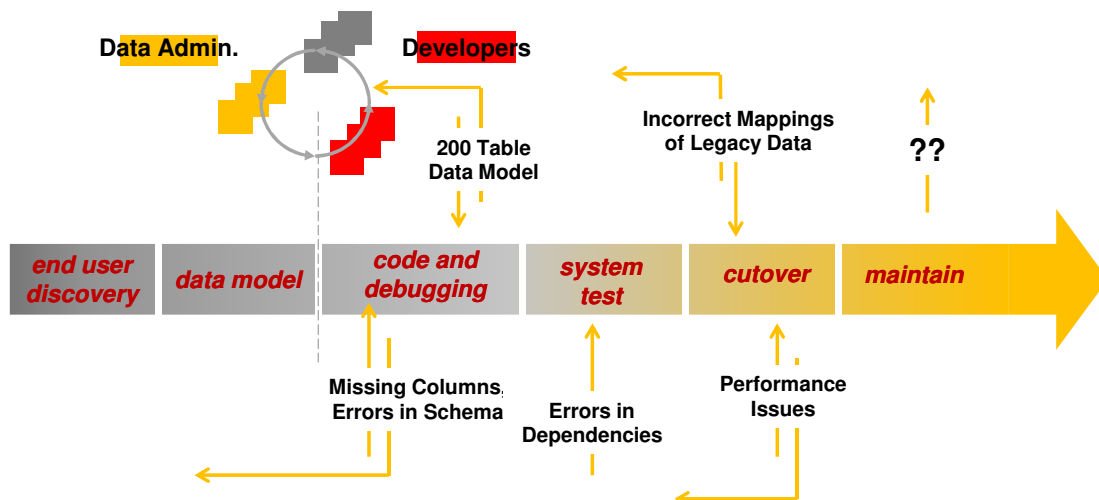
- Can I add X% more users ?
- Can I add this new application module ?
- Can I handle the expected/ increased seasonal service load/spike ?
- Can I downsize this system for cost ?
- What will happen to customer satisfaction if I do (YYY) while on line, or on peak ?
- Routine-Z is now slow; Did it grow slow over months, or overnight ?
(2 different problems/solutions)
- How can I improve application development velocity ?

2

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Legacy Application Development Lifecycle




3

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Creating a "Query Harness"-



Every statement; (SEL, INS/UPD, DEL, other)

- | | |
|---|-----------------|
| <ul style="list-style-type: none"> • Frequency • Concurrency • Spread (aka, Density) | Input |
| <ul style="list-style-type: none"> • Necessary SLA/SLO | Output/Measured |


4

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

NoSQLBench, Data-

Drastically reduce the time and effort to get sophisticated answers to simple questions-



- **Generated**
 - Less operational overhead (on-disk storage systems normally the limiting factor when performance testing)
 - Nothing to manage
- **Deterministic, statistically shaped, realistic data**
 - Example: First|Last|Full name binding based on USA survey data, Daniel more common than Melvin
 - (Same generated data on every execution)
 - 80+ built in bindings, each fully customizable
 - Built-in/custom weighting, skew
- **Out of the box workloads for,**
 - IOT, key/value, wide-column, UUID, other
 - Or, define your own
- **Drivers for,**
 - 11 systems out of the box; CQL, Http, Kafka, ..
 - Or, add your own
- **Built for distributed systems**
 - Partition aware operators

5


© 2020 Datastax, Inc. All rights reserved.

DATASTAX

NoSQLBench, (Service Layer)-

- **No programming-**
 - Single YAML file (a *Workload*), 3 or fewer common sections
 - (You can optionally design in JavaScript for very advanced needs)
- **Workload can be used variably, different;** *Scenarios*, scale, duration, targeting different/multiple nodes
- **Sync/Async, blocking/non-blocking**
 - Extremely simply configured
 - Extensible
- **Battle tested;** it's what DataStax uses internally for years, and at large scale (1 to hundreds of nodes)
- **Built in Prometheus/ Grafana container GUI** for high fidelity metrics, or 3 other routes to analyze
- **Easy to emulate your application (minutes of work)**

Drastically reduce the time and effort to get sophisticated answers to simple questions-

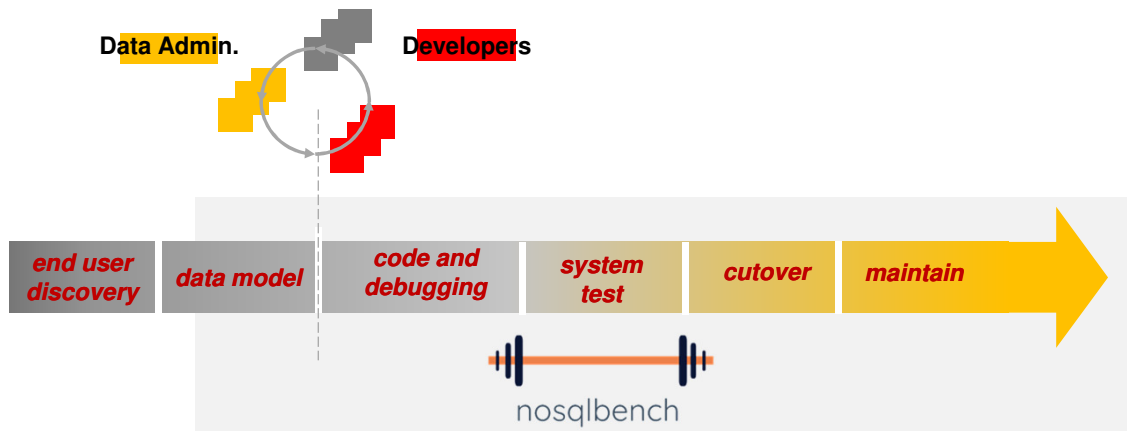


6

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Where/When to use NoSQLBench-



7

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Hands on Lab: Using a Built-in Workload, (Install verify, Hello World)

8

© 2020 Datastax, Inc. All rights reserved.



Hands on Lab: Overview (25 minutes or less to complete)

Pre-requisites-

- Linux system
- C*/DSE, or Pre-instantiated Astra instance

Here/now:

- Several simple LABS; *Learn-by-example*
- Get the 100-level; object hierarchy, use
- Which is faster; Primary Key, or SAI ?

Post-requisites- (not here/now, not formally)

- Dive into the GUI (Prometheus/Grafana)
- More database nodes; distributed Scenarios



9

© 2020 Datastax, Inc. All rights reserved.

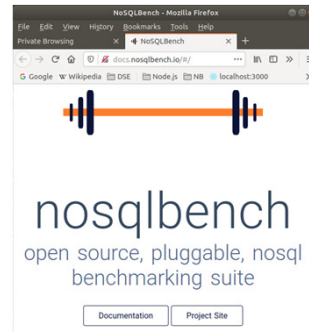


Step 01: Install, install-verify, (high level look)

docs.nosqlbench.io

<https://github.com/nosqlbench/nosqlbench/blob/main/DOWNLOADS.md>

- Get the binary (nb)
- `chmod 755 nb`
- From Linux Shell-
 - `nb --list-drivers` # That's two hyphens
 - `nb --list-workloads`
 - `nb --list-scenarios`
 - `nb --copy cql-keyvalue` # Overview / read this file
 - scenarios # Optional, used to variably (configure) a workload
 - bindings # the data generators
 - blocks
 - schema # the statements/activities proper, (these are just names, not keywords)
 - rampup # Generally DDL, run once, serially
 - main # Generally data creation, parallel
 - # The Query Harness proper



10

© 2020 Datastax, Inc. All rights reserved.



Step 02a: Copy, then edit, cql-iot.yaml, run (local, C*/DSE)

This step requires a C* or DSE operating locally and without authentication (without a password)

- `cp cql-keyvalue.yaml my_file.yaml`
- `vi my_file.yaml`
 - Under, scenarios → default # default is a keyword
 - Then both, rampup and main # not keywords, yes convention
 - Change the number (10 million) to just (10)
 - Exit and save
- From the command line, `nb my_file`
- Enter CQLSH, look for
 - A Keyspace titled, baselines
 - A Table titled, keyvalue
 - Count the rows-

11

© 2020 Datastax, Inc. All rights reserved.



Step 02b: Copy, then edit, cql-iot.yaml, run against Astra

This step requires an Astra account, and a download of your "secure connection bundle" Zip file

This Zip file contains your Keyspace/Database name, fyi, and that Database name is auto-generated into the file name. Below, ours is "my_database".

- `cp cql-keyvalue.yaml my_file.yaml`
- `vi my_file.yaml`
 - Under, scenarios → astra
 - Then both, rampup and main
 - Change the number (10 million) to just (10)
 - Exit and save
- From the command line,


```
nb my_file astra secureconnectbundle=secure-connect-my-database.zip
```
- From the Astra Home Page, enter CQLSH (or same basic steps, different means), look for
 - A keyspace titled, (your keyspace name)
 - A table titled, keyvalue
 - Count the rows-

12

© 2020 Datastax, Inc. All rights reserved.



Step 02: continued, *Optional*

- Using CQLSH, DROP the table


```
USE keyspace_name;           // Use correct value
DROP TABLE keyvalue;
```
- Add this clause to your "nb" command line,


```
tags:phase=schema
tags:phase=schema-astra    # if Astra
```
- Look around (check tables and data)
- Add one row of data,


```
INSERT INTO keyvalue (key, value)
VALUES ('xxx', 'yyy');
```
- Look around
- Add this clause to your "nb" command line,

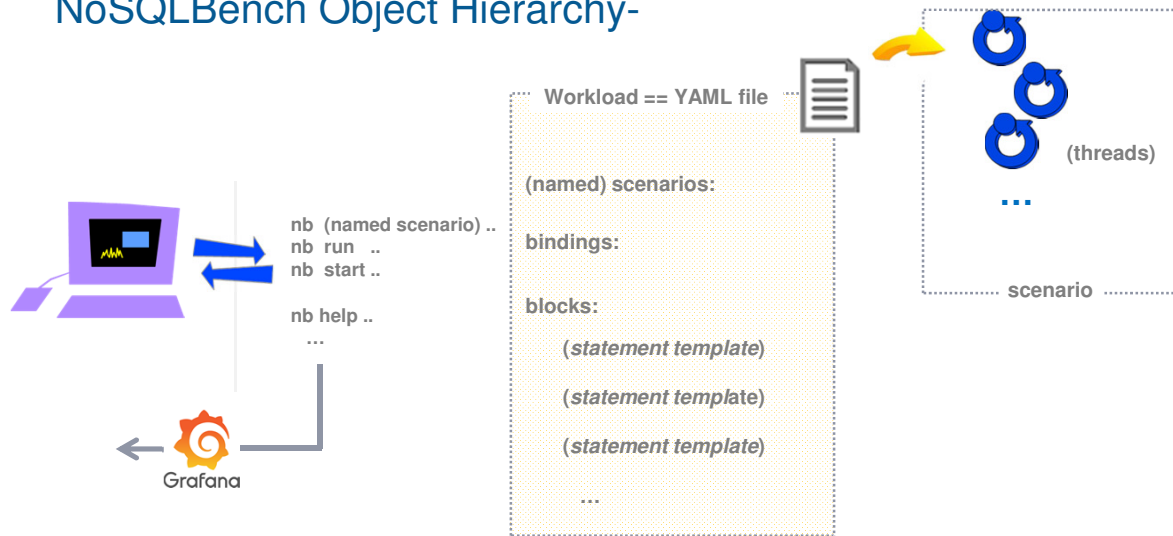

```
tags:phase=rampup
```
- Look around

13

© 2020 Datastax, Inc. All rights reserved.



NoSQLBench Object Hierarchy-



14

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

NoSQLBench Object Hierarchy: Scenarios

description: A workload with only text keys and text values

scenarios:

default:

```
schema: run driver=cql tags==phase:schema threads==1 cycles==UNDEF
rampup: run driver=cql tags==phase:rampup cycles===TEMPLATE(rampup-cycles,10) threads=auto
main: run driver=cql tags==phase:main cycles===TEMPLATE(main-cycles,10) threads=auto
astra:
schema: run driver=cql tags==phase:schema-astra threads==1 cycles==UNDEF
rampup: run driver=cql tags==phase:rampup cycles===TEMPLATE(rampup-cycles,10) threads=auto
main: run driver=cql tags==phase:main cycles===TEMPLATE(main-cycles,10) threads=auto
```

- (Named) scenarios
- 'default'
- run|start
- driver
- 1, 2, or 3 equal signs
- tags
- threads
 - 1
 - (Default, auto)
- cycles
 - UNDEF
 - 10 aka, 0..10
- (others, not shown)

15

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

NoSQLBench Object Hierarchy: Bindings

- Many dozens to choose from
 - Different (types); deterministic, random, other
 - name : recipe
- Top level, also lower scope
- <<label:default>>
 - Not just for bindings, available elsewhere
 - Override passed on command line
- Casting too, E.g., .. -> int

```
bindings:
  seq_key: Mod(<<keycount:1000000000>>); ToString() -> String
  seq_value: Hash(); Mod(<<valuecount:1000000000>>); ToString() -> String
  rw_key: <<keydist:Uniform(0,1000000000)->int>>; ToString() -> String
  rw_value: Hash(); <<valdist:Uniform(0,1000000000)->int>>; ToString() -> String
```

16

© 2020 Datastax, Inc. All rights reserved.



NoSQLBench Object Hierarchy: Statement (Templates)

```
- name: schema-astra
  tags:
    phase: schema-astra
  params:
    prepared: false
  statements:
    - create-table: |
        create table if not exists <<keyspace:baselines>>.<<table:keyvalue>> (
          key text,
          value text,
          PRIMARY KEY (key)
        );
      tags:
        name: create-table-astra
- name: rampup
  tags:
    phase: rampup
  params:
    cl: <<write_cl:LOCAL_QUORUM>>
  statements:
    - rampup-insert: |
        insert into <<keyspace:baselines>>.<<table:keyvalue>>
          (key, value)
        values ({seq_key},{seq_value});
      tags:
        name: rampup-insert
```

- Can have multiple statements per phase
- CQL DDL
 - Not prepared
 - **threads==1**
 - **cycles == UNDEF**
- Tags/name, execute subset
- **driver=stdout**
- <<write_cl: .. >>
- {seq_key}

17

© 2020 Datastax, Inc. All rights reserved.



Step 03: Checkpoint



- Why is it a better idea to develop against data dense nodes ?
- 400 concurrent users execute a given routine with 3 SELECTS and 1 INSERT every 60 seconds, with a 20 second think-time/pause-
 - What are the Frequency, Concurrency, and Spread ?
 - What is the target SLA/SLO ?
- In NoSQLBench parlance, the YAML file is termed a _____ ?
- What are the 3 most common sections in that YAML ? Describe each ?
- What's the difference between; =, ==, === ?
- What does a << k : v >> represent ?
- What does a { bb } represent ?

18

© 2020 Datastax, Inc. All rights reserved.



Hands on Lab: scenarios, schema phase

19

© 2020 Datastax, Inc. All rights reserved.



Step 04: Design/Run your own Workload (your own YAML file)

- Assuming you know CQL, 98% of your errors will be indentation in the YAML- Refer-to/copy from our earlier YAML
- driver=stdout override later to CQL
- schema phase
 - Drop a keyspace
 - Create a keyspace
 - Create a table; 10 columns (really, we'll use 2)
 - Single column PK, type TEXT
 - Addition column, type TEXT, will contain same value as the PK
 - CREATE an SAI index



Not conclusive because single node, low power nodes, but;

How close is a Primary Key lookup to a secondary index lookup ??

20

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 04: Design/Run our own Workload, sample CQL



```
DROP KEYSPACE IF EXISTS ks_44;
CREATE KEYSPACE ks_44
  WITH replication = {'class': 'SimpleStrategy',
    'replication_factor': '1'};
```

C*/DSE only

```
CREATE TABLE ks_44.t1
(
  col1 TEXT PRIMARY KEY,
  col4 TEXT
);
CREATE CUSTOM INDEX col4_idx
  ON ks_44.t1 (col4) USING 'StorageAttachedIndex';

INSERT INTO ks_44.t1 (col1, col4) VALUES ('eee', 'fff');
SELECT * FROM ks_44.t1 WHERE col4 = 'fff';
```

21

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 04: Get the schema phase working



As stated:

- Get the schema phase working; keyspace, table, index
- driver=stdout then, driver=cql
- Errors with stdout, it's your YAML indents
- Errors with cql, copy and paste into CQLSH, debug that way
- Cheating; answer on next few pages

22

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 05: Checkpoint, answer so far-

scenarios:

default:

```
schema: run driver=stdout tags==phase:schema threads==1 cycles==UNDEF
# more stuff goes here later
```

bindings:

```
# stuff will go here later
```

blocks:

- tags:

```
phase: schema
```

params:

```
prepared: false
```

statements:



23

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 05: Checkpoint, answer so far-

```
- drop_keyspace: |
  DROP KEYSPACE IF EXISTS <<keyspace:ks_44>>;
tags:
  name: drop_keyspace

- create_keyspace: |
  CREATE KEYSPACE <<keyspace:ks_44>>
  WITH replication = {'class': 'SimpleStrategy',
    'replication_factor': '1'};
tags:
  name: create_keyspace
```



24

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 05: Checkpoint, answer so far-

```
- create_table: |
  CREATE TABLE <<keyspace:ks_44>>.<<table:t1>>
  (
    col1 TEXT PRIMARY KEY,
    col2 TEXT,
    col3 TEXT,
    col4 TEXT,
    col5 TEXT,
    col6 TEXT,
    col7 TEXT,
    col8 TEXT,
    col9 TEXT,
    col10 TEXT
  );
tags:
  name: create_table
```



25

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 05: Checkpoint, answer so far-

```
- create_index: |
  CREATE CUSTOM INDEX col4_idx
    ON <<keyspace:ks_44>>.<<table:t1>> (col4) USING 'StorageAttachedIndex';
tags:
  name: create_index
```



26

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Step 05: Checkpoint, answer so far-

To execute on the command line,

```
nb (file_name)
nb (file_name) driver=cql
nb (file_name) driver=cql secureconnectionbundle= ...
```

From the answer provided so far; What else needs to change for Astra ?



27

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Hands on Lab: rampup phase

28

© 2020 Datastax, Inc. All rights reserved.



Step 06: Add rampup phase

- 10 rows only until absolute end of exercise, until you are positive everything is perfect
- Why ? Even those DROP KEYSPACE statements are not (entirely) free)
- The primary key column and col4 should use the same binding (receive the same generated value)
- Experiment with other bindings for the other columns
- *In addition to more statements, you need to add to scenarios, and bindings*



29

© 2020 Datastax, Inc. All rights reserved.



Step 06: Aids/snippets (these are not complete)

```
nb run workload=(file_name) driver=stdout tags=phase:rampup cycles=10
```

bindings:

```
# there are much easier solutions to the first line below; we prefer this one
colx: Mod(<<uuidCount:100000000>>L); ToHashedUUID() -> java.util.UUID; ToString() -> String
```

```
col8: FullNames() -> String
```

- tags:

```
phase: rampup
params:
prepared: true
statements:
```

```
- insert: |
  INSERT INTO <<keyspace:ks_44>>.<<table:t1>> (col1, col4, col8)
  VALUES ( {colx}, {colx}, {col8} );
tags:
name: insert
```

30

© 2020 Datastax, Inc. All rights reserved.



Step 07: Checkpoint



- Using CQLSH, check for 10 rows in your table
- Could FullNames() have been used as the PK ?
- What did the new/sample "nb" command line invocation do differently ?

31

© 2020 Datastax, Inc. All rights reserved.



Hands on Lab: free Astra tier specific notes

32

© 2020 Datastax, Inc. All rights reserved.



Notes specific to the free Astra tier

- This tier is free, *and* tiny
- INSERTing 100 rows/sec is super safe, higher than that and ..
- You'll likely need to add,



```
nb run workload=(file_name) driver=stdout tags=phase:rampup cycles=10 cyclerate=1
```

*Above calls to run only one statement per second (for test)
You can do 100 on the Astra free tier, maybe as high as 600*

33

© 2020 Datastax, Inc. All rights reserved.



Hands on Lab: main phase

34

© 2020 Datastax, Inc. All rights reserved.



Step 08: Design/Run our own Workload, main phase

- Still .. 10 rows only, until you're certain it's perfect
- Yes, any test under (n) million rows may have accuracy concerns
- **SELECT using the PK**
- **SELECT using an SAI indexed column (col4)**
- **View results**

35

© 2020 Datastax, Inc. All rights reserved.



Step 08: Aids/snippets (these are not complete)

nb run workload=(file_name) driver=stdout tags=phase:main cycles=10

nb run workload=sol3 driver=cql tags=name:query1 cycles=10

(No new bindings needed, *although we could even further randomize the lookup key*)

```
- tags:
  phase: main
  params:
    prepared: true
  statements:

- query1: |
  SELECT * FROM <<keyspace:ks_44>>.<<table:t1>> WHERE col1 = {colx} ;
  tags:
    name: query1

- query1: |
  SELECT * FROM <<keyspace:ks_44>>.<<table:t1>> WHERE col4 = {colx} ;
  tags:
    name: query2
```

36

© 2020 Datastax, Inc. All rights reserved.



Step 08: Aids/snippets (these are not complete)

- Several paths to view results. One is to view the 'logs', created by default, below the current working directory

```
• rm -r logs
• (Re-run just 'query1', then 'query2' to driver=cql
• cd logs
• vi scenario*.log      # this file name is output when you run "nb"
```

- Less than 100 lines, probably; see next page

- And Yes, there's also Grafana, we'll *introduce* that shortly

37

© 2020 Datastax, Inc. All rights reserved.



Step 08: Viewing the (scenario results) log file

- "Results" is its own whole set of topics-
 - Start here: http://docs.nosqlbench.io/#/docs/getting_started/03_reading_metrics
- Lots to look at, but .. let's look at,

```
2020-08-12 13:44:48,293 INFO [main] i.n.e.c.ScenarioResult [Slf4jReporter.java:374]
type=TIMER, name=sol3.result-success, count=10, min=1046.496, max=5771.263,
mean=2666.6912, stddev=1445.067675240492, median=2332.415, p75=3398.399, p95=5771.263,
p98=5771.263, p99=5771.263, p999=5771.263, mean_rate=18.600062259244396, m1=0.0, m5=0.0,
m15=0.0, rate_unit=events/second, duration_unit=microseconds
```

```
2020-08-12 13:44:47,957 INFO [main] i.n.e.c.ScenarioResult [Slf4jReporter.java:374]
type=HISTOGRAM, name=sol3.resultset-size, count=10, min=0, max=1, mean=0.1,
stddev=0.30000000000000004, median=0.0, p75=0.0, p95=1.0, p98=1.0, p99=1.0, p999=1.0
```

38

© 2020 Datastax, Inc. All rights reserved.



Step 08: last page was 10 rows, the PK lookup, now SAI

```
name=sol3.result-success, count=10, min=1046.496, max=5771.263, mean=2666.6912
name=sol3.resultset-size, count=10, min=0, max=1
```

```
name=sol3.result-success, count=10, min=1222.784, max=5263.615, mean=2126.1504
name=sol3.resultset-size, count=10, min=0, max=1
```



What should we change ?

39

© 2020 Datastax, Inc. All rights reserved.



Step 08: 10 rows, then 10M rows

10 Rows, 10 queries

`name=sol3.result-success, count=10, min=1046.496, max=5771.263, mean=2666.6912` PK lookup
`name=sol3.resultset-size, count=10, min=0, max=1`
`name=sol3.result-success, count=10, min=1222.784, max=5263.615, mean=2126.1504` SAI lookup
`name=sol3.resultset-size, count=10, min=0, max=1`

`name=sol3.result-success, count=10000, min=123.524, max=10125.311, mean=458.1033824` PK lookup
`name=sol3.resultset-size, count=10000, min=1, max=1`
`name=sol3.result-success, count=10000, min=269.552, max=78245.887, mean=1546.370848` SAI lookup
`name=sol3.resultset-size, count=10000, min=1, max=1`

10M Rows, 1000 queries

40

© 2020 Datastax, Inc. All rights reserved.



Step 09: If time allows

- Prometheus/Grafana
- Add to command line first time,
--docker-metrics

Different switch on subsequent use



41

© 2020 Datastax, Inc. All rights reserved.



Further Study (next steps):

- **More on the process model-**
 - **run|start** Blocking, non-blocking
 - **Sync/Async**
 - **More control over the (statements)**
 - **Multiple "nb" hosts** (simple, but not covered..)
- **More on parameters**
- **More on bindings;** which are deterministic, random
- **More on the object hierarchy**
- **Scripting**
- **More on results;** client side (nb), server side (not nb)
- **Basically a lot more**



42

© 2020 Datastax, Inc. All rights reserved.

DATASTAX



Thank You!

Final answer: Complete solution

```
# Run via,
#
# nb (file_name)
# nb (file_name) driver=cql
#
# nb run workload=(file_name) driver=stdout tags=phase:rampup cycles=10
#
# nb run workload=(file_name) driver=stdout tags=name:query1 cycles=10

# Because of the DROP/CREATE KEYSPACE, this file does not run against Astra
```

scenarios:

default:

```
schema: run driver=stdout tags==phase:schema threads==1 cycles==UNDEF
rampup: run driver=stdout tags==phase:rampup threads==auto cycles=10000000
main: run driver=stdout tags==phase:main threads==auto cycles=100000
```

bindings:

```
colx: Mod(<<uuidCount:10000000>>L); ToHashedUUID() -> java.util.UUID; ToString() -> String
col8: FullNames() -> String
```



44

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Final answer: Complete solution

blocks:

- tags:

phase: schema

params:

prepared: false

statements:

- drop_keyspace: |

```
DROP KEYSPACE IF EXISTS <<keyspace:ks_44>>;
```

tags:

name: drop_keyspace

- create_keyspace: |

```
CREATE KEYSPACE <<keyspace:ks_44>>
```

```
WITH replication = {'class': 'SimpleStrategy',
```

```
'replication_factor': '1'};
```

tags:

name: create_keyspace



45

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Final answer: Complete solution

```
- create_table: |
  CREATE TABLE <<keyspace:ks_44>>.<<table:t1>>
  (
    col1 TEXT PRIMARY KEY,
    col2 TEXT,
    col3 TEXT,
    col4 TEXT,
    col5 TEXT,
    col6 TEXT,
    col7 TEXT,
    col8 TEXT,
    col9 TEXT,
    col10 TEXT
  );
tags:
  name: create_table

- create_index: |
  CREATE CUSTOM INDEX col4_idx
  ON <<keyspace:ks_44>>.<<table:t1>> (col4) USING 'StorageAttachedIndex';
tags:
  name: create_index
```



46

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Final answer: Complete solution

```
- tags:
  phase: rampup
params:
  prepared: true
statements:

- insert: |
  INSERT INTO <<keyspace:ks_44>>.<<table:t1>> (col1, col4, col8)
  VALUES ( {colx}, {colx}, {col8} );
tags:
  name: insert
```



47

© 2020 Datastax, Inc. All rights reserved.

DATASTAX

Final answer: Complete solution

```
- tags:  
  phase: main  
params:  
  prepared: true  
statements:  
  
- query1: |  
  SELECT * FROM <<keyspace:ks_44>>.<<table:t1>> WHERE col1 = {colx} ;  
tags:  
  name: query1  
  
- query1: |  
  SELECT * FROM <<keyspace:ks_44>>.<<table:t1>> WHERE col4 = {colx} ;  
tags:  
  name: query2
```

